

# IMDBClassifier

December 14, 2021

```
[1]: #Implementing the IMDB movie review classifier from Chapter 3.4 in Deep Learning with Python
```

```
#Importing the necessary libraries
from keras.datasets import imdb
import numpy as np
from keras import models
from keras import layers
import matplotlib.pyplot as plt
```

```
[2]: #Importing and splitting the data into train and test data
(train_data, train_labels), (test_data, test_labels) = imdb.load_data(num_words=10000)
```

```
#Looking at the data shape
print(train_data.shape)
print(test_data.shape)
```

(25000,)

(25000,)

```
[3]: #One-hot encoding the training and testing data
def vectorize_sequences(sequences, dimension=10000):
    results = np.zeros((len(sequences), dimension))
    for i, sequence in enumerate(sequences):
        results[i, sequence] = 1.
    return results
```

```
x_train = vectorize_sequences(train_data)
x_test = vectorize_sequences(test_data)
y_train = np.asarray(train_labels).astype('float32')
y_test = np.asarray(test_labels).astype('float32')
```

```
[4]: #Defining the layers of the neural network
model = models.Sequential()
model.add(layers.Dense(16, activation = 'relu', input_shape = (10000,)))
model.add(layers.Dense(16, activation = 'relu'))
```

```
model.add(layers.Dense(1, activation = 'sigmoid'))
```

```
[5]: #Compiling the model and passing the optimizer, loss function, and required
      ↪metrics into the model
model.compile(optimizer = 'rmsprop', loss = 'binary_crossentropy', metrics =
      ↪['accuracy'])
```

```
[6]: #Setting aside a validation dataset
x_val = x_train[:10000]
partial_x_train = x_train[10000:]
y_val = y_train[:10000]
partial_y_train = y_train[10000:]
```

```
[7]: #Training the model with 20 epochs, a batch size of 512, and validating with
      ↪the validation data
history = model.fit(partial_x_train,
                    partial_y_train,
                    epochs = 20,
                    batch_size = 512,
                    validation_data = (x_val, y_val))
```

Epoch 1/20

30/30 [=====] - 1s 39ms/step - loss: 0.5237 - accuracy: 0.7725 - val\_loss: 0.3811 - val\_accuracy: 0.8702

Epoch 2/20

30/30 [=====] - 1s 38ms/step - loss: 0.3043 - accuracy: 0.9019 - val\_loss: 0.2987 - val\_accuracy: 0.8882

Epoch 3/20

30/30 [=====] - 2s 51ms/step - loss: 0.2186 - accuracy: 0.9295 - val\_loss: 0.2804 - val\_accuracy: 0.8874

Epoch 4/20

30/30 [=====] - 1s 32ms/step - loss: 0.1708 - accuracy: 0.9459 - val\_loss: 0.2917 - val\_accuracy: 0.8825

Epoch 5/20

30/30 [=====] - 1s 32ms/step - loss: 0.1401 - accuracy: 0.9545 - val\_loss: 0.2902 - val\_accuracy: 0.8862

Epoch 6/20

30/30 [=====] - 1s 26ms/step - loss: 0.1146 - accuracy: 0.9629 - val\_loss: 0.3179 - val\_accuracy: 0.8822

Epoch 7/20

30/30 [=====] - 1s 27ms/step - loss: 0.0897 - accuracy: 0.9743 - val\_loss: 0.3210 - val\_accuracy: 0.8835

Epoch 8/20

30/30 [=====] - 1s 27ms/step - loss: 0.0743 - accuracy: 0.9801 - val\_loss: 0.3353 - val\_accuracy: 0.8823

Epoch 9/20

30/30 [=====] - 1s 27ms/step - loss: 0.0611 - accuracy:

```

0.9833 - val_loss: 0.3700 - val_accuracy: 0.8800
Epoch 10/20
30/30 [=====] - 1s 27ms/step - loss: 0.0514 - accuracy:
0.9871 - val_loss: 0.3848 - val_accuracy: 0.8794
Epoch 11/20
30/30 [=====] - 1s 30ms/step - loss: 0.0400 - accuracy:
0.9908 - val_loss: 0.4275 - val_accuracy: 0.8771
Epoch 12/20
30/30 [=====] - 1s 28ms/step - loss: 0.0286 - accuracy:
0.9945 - val_loss: 0.4498 - val_accuracy: 0.8725
Epoch 13/20
30/30 [=====] - 1s 38ms/step - loss: 0.0245 - accuracy:
0.9953 - val_loss: 0.4900 - val_accuracy: 0.8738
Epoch 14/20
30/30 [=====] - 2s 54ms/step - loss: 0.0193 - accuracy:
0.9965 - val_loss: 0.5133 - val_accuracy: 0.8705
Epoch 15/20
30/30 [=====] - 1s 35ms/step - loss: 0.0135 - accuracy:
0.9986 - val_loss: 0.5606 - val_accuracy: 0.8664
Epoch 16/20
30/30 [=====] - 1s 27ms/step - loss: 0.0100 - accuracy:
0.9996 - val_loss: 0.7403 - val_accuracy: 0.8511
Epoch 17/20
30/30 [=====] - 1s 31ms/step - loss: 0.0084 - accuracy:
0.9988 - val_loss: 0.6209 - val_accuracy: 0.8681
Epoch 18/20
30/30 [=====] - 1s 31ms/step - loss: 0.0103 - accuracy:
0.9979 - val_loss: 0.6545 - val_accuracy: 0.8674
Epoch 19/20
30/30 [=====] - 1s 28ms/step - loss: 0.0035 - accuracy:
0.9999 - val_loss: 0.6872 - val_accuracy: 0.8671
Epoch 20/20
30/30 [=====] - 1s 30ms/step - loss: 0.0064 - accuracy:
0.9987 - val_loss: 0.7188 - val_accuracy: 0.8664

```

```

[8]: #Obtaining the training history and viewing its keys
history_dict = history.history
history_dict.keys()

```

```

[8]: dict_keys(['loss', 'accuracy', 'val_loss', 'val_accuracy'])

```


```

[9]: #Plotting the training loss and the validation loss for each epoch

loss_values = history_dict['loss']
val_loss_values = history_dict['val_loss']
epochs = range(1, len(loss_values) + 1)

```

```
plt.plot(epochs, loss_values, 'bo', label = 'Training loss')
plt.plot(epochs, val_loss_values, 'b', label = 'Validation loss')
plt.title('Training and validation loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()
plt.show()
```




[10]: *#Plotting the training accuracy and the validation accuracy for each epoch*

```
plt.clf()

acc_values = history_dict['accuracy']
val_acc_values = history_dict['val_accuracy']

plt.plot(epochs, acc_values, 'bo', label = 'Training accuracy')
plt.plot(epochs, val_acc_values, 'b', label = 'Validation accuracy')
plt.title('Training and validation accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()
plt.show()
```



```
[11]: #Evaluating the model on the test dataset
results = model.evaluate(x_test, y_test)
print('The loss is {} ({} epochs)'.format(round(results[0], 4), len(epochs)))
print('The accuracy is {} percent ({} epochs)'.format(round(results[1], 4)*100, len(epochs)))
```

```
782/782 [=====] - 2s 2ms/step - loss: 0.7851 -
accuracy: 0.8513
The loss is 0.7851 (20 epochs)
The accuracy is 85.13 percent (20 epochs)
```

```
[12]: #Retraining the model, this time with 4 epochs, and viewing its loss and
      ↪ accuracy

#Defining the layers
model = models.Sequential()
model.add(layers.Dense(16, activation = 'relu', input_shape = (10000,)))
model.add(layers.Dense(16, activation = 'relu'))
model.add(layers.Dense(1, activation='sigmoid'))

#Defining the optimizer, loss function, and metrics
model.compile(optimizer = 'rmsprop', loss = 'binary_crossentropy', metrics = [
      ↪ 'accuracy'])
```

```
#Fitting the model with the training data, 4 epochs, and a batch size of 512
model.fit(x_train, y_train, epochs = 4, batch_size = 512)
```

```
Epoch 1/4
49/49 [=====] - 0s 9ms/step - loss: 0.4490 - accuracy:
0.8247
Epoch 2/4
49/49 [=====] - 0s 9ms/step - loss: 0.2608 - accuracy:
0.9100
Epoch 3/4
49/49 [=====] - 0s 8ms/step - loss: 0.2016 - accuracy:
0.9288
Epoch 4/4
49/49 [=====] - 0s 7ms/step - loss: 0.1698 - accuracy:
0.9394
```

```
[12]: <tensorflow.python.keras.callbacks.History at 0x7faf03ff2490>
```

```
[13]: #Obtaining the results
results = model.evaluate(x_test, y_test)
print('The loss is {} ({} epochs)'.format(round(results[0], 4), 4))
print('The accuracy is {} percent ({} epochs)'.format(round(results[1], 4)*100, 4))
↪4))
```

```
782/782 [=====] - 2s 2ms/step - loss: 0.2910 -
accuracy: 0.8839
The loss is 0.291 (4 epochs)
The accuracy is 88.39 percent (4 epochs)
```

```
[14]: #Generating predictions using the test data
#Generating the likelihood of movie reviews being positive
model.predict(x_test)
```

```
[14]: array([[0.20587534],
             [0.99967194],
             [0.88956904],
             ...,
             [0.13893089],
             [0.07124329],
             [0.6224241 ]], dtype=float32)
```